

Autograd, optimizare

Bibliografie:

1. Deep Learning with PyTorch, Eli Stevens, Luca Antiga, Thomas Viehmann, Manning; 1st edition (August 4, 2020)
2. PyTorch Pocket Reference: Building and Deploying Deep Learning Models, by Joe Papa, O'Reilly Media; 1st edition (June 1, 2021)

```
In [1]: import torch  
print(f'Versiune torch: {torch.__version__}')
```

Versiune torch: 2.0.0+cu118

```
In [2]: torch.cuda.is_available()
```

Out[2]: True

Calcul de gradienti cu autograd

Calculul derivatelor parțiale ale unei funcții este crucial pentru algoritmi actuali de machine learning. Modificând ponderile (parametrii, coeficienții) unei funcții, se poate face printr-un proces iterativ aducerea funcției într-o zonă în care gradientii să fie 0.

Exemplul 1: funcție de o singură variabilă

```
In [3]: # definim o variabilă x. Pentru ca se va calcula derivata unei funcții în raport cu x,  
# faptul că trebuie considerat x pentru calcul de gradient  
x = torch.tensor(10.0, requires_grad = True)  
print(x)  
# x trebuie dat ca valoare floating point;  
# o valoare inițială dată ca întreg (10) face ca tensorul să aibă datatype întreg.  
# În PyTorch gradientii se pot calcula numai pentru cantități floating point
```

tensor(10., requires_grad=True)

Verificăm dacă tensorul x are gradient definit în acest moment:

```
In [4]: x.grad is None
```

Out[4]: True

Definim o expresie care depinde de x , de exemplu $y = 5x^3 - x^2 + 2x - 100$. Derivata lui y în funcție de x este $\frac{dy}{dx} = 15x^2 - 2x + 2$ care evaluată pentru $x = 10$ da rezultatul 1482.

```
In [5]: y = 5*x**3 - x**2 + 2*x - 100  
# declansam calculul de gradient. Metoda backward() poate fi chemată pentru un tensor
```

```
y.backward()
```

Din acest moment, tensorul x are gradient, calculat de metoda `backward` :

```
In [6]: # verificam daca s-a calculat un gradient pentru x, corespunzator lui dy/dx:
print(f'Avem gradient pe x: {x.grad is not None}')
```

Avem gradient pe x: True

```
In [7]: # recuperam gradientul:
dydx = x.grad
dydx
```

```
Out[7]: tensor(1482.)
```

Important! Daca mai apelam odata `backward()` pe o expresie care il foloseste pe x (de exemplu: aceeasi expresie y ca mai sus), atunci se calculeaza din nou gradient pentru vectorul x si se acumuleaza la cel calculat deja. Ca atare, pentru acelasi y noul gradient este $1482 + 1482 = 2964$:

```
In [8]: y = 5*x**3 - x**2 + 2*x - 100
y.backward()
x.grad
```

```
Out[8]: tensor(2964.)
```

Exemplul 2: functie de 2 variabile

Consideram functia: $z(x, y) = x^2 + y^5 - x^2 \cdot y$ Derivatele partiale sunt: $\frac{\partial z}{\partial x} = 2x - 2xy$ respectiv $\frac{\partial z}{\partial y} = 5y^4 - x^2$

Pentru $x=3, y=4$ cele doua derivate partiale au valorile: $\frac{\partial z}{\partial x} = -18$ si respectiv $\frac{\partial z}{\partial y} = 1271$

```
In [9]: x=torch.tensor(3.0, requires_grad=True)
y=torch.tensor(4.0, requires_grad=True)
z=x**2+y**5-x**2*y
z.backward()
```

```
In [10]: # recuparam gradientii dz/dx, dz/dy
print(f'gradient dupa x={x.grad}\ngradient dupa y={y.grad}')
```

```
gradient dupa x=-18.0
gradient dupa y=1271.0
```

Exemplul 3: functie de eroare patratica, mai multe ponderi instruibile

In PyTorch se pot calcula automat gradientii chiar daca sunt folositi in interiorul unor functii definite de utilizator.

```
In [11]: def f(w: torch.Tensor, x: torch.Tensor) -> torch.Tensor:
          return x.dot(w)

          def error_function(y: torch.Tensor, y_hat: torch.Tensor) -> torch.Tensor:
              return (y - y_hat) ** 2

          # x: valori de intrare
          x = torch.tensor([1.0, 4.0, -3.0, 5.0, 4.0])
          print(f'Verificam daca tensorul x e pregatit sa aiba gradient calculat: {x.requires_grad_}')
          # w: ponderi
          w = torch.rand_like(x, requires_grad=True)

          ground_truth = 10

          # eroare patratica medie
          error = error_function(ground_truth, f(w, x))
          print(f'Eroare de predictie: error={error}')

          # se calculeaza gradientii de catre Pytorch
          error.backward()
          print(f'Gradientii pentru w: {w.grad}')

          print(f'Gradientii pentru x nu se calculeaza, deoarece tensorul x are requires_grad=False')
          print(f'\tx.grad is None={x.grad is None}')
```

Verificam daca tensorul x e pregatit sa aiba gradient calculat: False
 Eroare de predictie: error=23.257722854614258
 Gradientii pentru w: tensor([-9.6453, -38.5810, 28.9358, -48.2263, -38.5810])
 Gradientii pentru x nu se calculeaza, deoarece tensorul x are requires_grad=False
 x.grad is None=True

```
In [12]: # aplicam un gradient descent manual: din ponderile actuale se scade gradientul
          # inmultit cu un learning rate mic
          lr = 0.01
          w = w - lr * w.grad
          error = error_function(ground_truth, f(w, x))
          print(f'Noua eroare: {error}')
          # se constata scaderea erorii
```

Noua eroare: 2.6885907649993896

Optimizare manuala

```
In [13]: # model de predictie
          def model(x, w, b):
              return w * x + b

          # functia de cost
          def error(y_hat: torch.Tensor, y: torch.Tensor, weights: torch.Tensor, lambda: float):
              quality_error = torch.mean((y_hat - y)**2)
              regularization = lambda * weights.norm() ** 2
              return quality_error + regularization
```

```
In [15]: import matplotlib.pyplot as plt

          params = torch.tensor([1.0, 0.0], requires_grad=True)
          x = torch.tensor(20.0)
```

```
# la inceput, gradientii pt ponderi nu sunt calculati
print(f'params.grad is None: {params.grad is None}')

y_true = torch.tensor(10)
learning_rate = 0.0001
n_iters = 100
lmbda = 1.0

losses = []

for i in range(1, n_iters+1):
    # gradientii calculati in mod repetat se acumuleaza (insumeaza cu gradientii calculati)
    # evitam aceasta insumare prin setarea lor explicita la 0
    if params.grad is not None:
        params.grad.zero_()

    y_hat = model(x, *params)
    loss = error(y_hat, y_true, params[1:], lmbda)
    losses.append(loss.detach().item())
    if i % 10 == 0:
        print(f'Iteration: {i}, loss: {loss}')
    loss.backward()

    with torch.no_grad():
        assert params.grad is not None
        params -= learning_rate * params.grad
```


```
params.grad is None: True
Iteration: 10, loss: 22.206819534301758
Iteration: 20, loss: 4.172554969787598
Iteration: 30, loss: 0.7844100594520569
Iteration: 40, loss: 0.1478678286075592
Iteration: 50, loss: 0.028277603909373283
Iteration: 60, loss: 0.005807911511510611
Iteration: 70, loss: 0.001584485056810081
Iteration: 80, loss: 0.0007890488486737013
Iteration: 90, loss: 0.0006376549135893583
Iteration: 100, loss: 0.0006072467076592147
```

Nota: daca celula de plot produce eroare, atunci se foloseste workaround de la

<https://stackoverflow.com/questions/53014306/error-15-initializing-libiomp5-dylib-but-found-libiomp5-dylib-already-initial>:

```
In [16]: import os
os.environ["KMP_DUPLICATE_LIB_OK"]="TRUE"
```

```
In [17]: plt.plot(losses)
plt.xlabel('Epoca')
plt.ylabel('Cost')
plt.show()
```



Optimizatori Pytorch

Mai sus s-a implementat in mod manual algoritmul de SGD, doar pentru a demonstra calculul gradientilor si efectul asupra ponderilor. De regula, se foloseste un agloritm de optimizare predefinit care aplica gradientii pe ponderi. O lista a optimizatorilor implementati in Pytorch este data mai jos.

```
In [18]: import torch.optim as optim  
[opt for opt in dir(optim) if opt[0].isupper()]
```

```
Out[18]: ['ASGD',  
          'Adadelata',  
          'Adagrad',  
          'Adam',  
          'AdamW',  
          'Adamax',  
          'LBFGS',  
          'NAdam',  
          'Optimizer',  
          'RAdam',  
          'RMSprop',  
          'Rprop',  
          'SGD',  
          'SparseAdam']
```

Fiecare optimizator preia o lista de parametri (de ex ponderile unui model de predictie). Dupa ce gradientii pentru minibatchul curent sunt calculati, se aplica in mod corespunzator pe ponderi.

```
In [19]: x = torch.tensor(20.0)
         params = torch.tensor([1.0, 0.0], requires_grad=True)
         learning_rate = 1e-3
         optimizer = optim.SGD([params], lr=learning_rate)
```

Secventa care urmeaza are acelasi scop ca si ciclarea din sectiunea anterioara: se modifica ponderile pentru a face minimizarea functiei `error`.

```
In [20]: losses = []

         for epoch in range(1, n_iters + 1):
             y_hat = model(x, *params)
             loss = error(y_hat, y_true, params[1:], lambda)
             losses.append(loss.item())
             print(f'Iteration: {epoch}, loss: {loss}')
             # se pun gradientii anterior calculati pe 0. Intrucat optimizatorul este responsabil
             # managementul ponderilor, resetarea gradientilor este facuta prin intermediul lui
             optimizer.zero_grad()


             # calcul de derivate partiale
             loss.backward()

             # se modifica ponderile folosind gradientii calculati de apelul backward()
             optimizer.step()
```

Iteration: 1, loss: 100.0
Iteration: 2, loss: 3.9207983016967773
Iteration: 3, loss: 0.15429987013339996
Iteration: 4, loss: 0.006642207968980074
Iteration: 5, loss: 0.0008512443164363503
Iteration: 6, loss: 0.000621881103143096
Iteration: 7, loss: 0.0006105515058152378
Iteration: 8, loss: 0.000607771216891706
Iteration: 9, loss: 0.0006053348770365119
Iteration: 10, loss: 0.0006029214127920568
Iteration: 11, loss: 0.0006005181348882616
Iteration: 12, loss: 0.000598124461248517
Iteration: 13, loss: 0.0005957404500804842
Iteration: 14, loss: 0.0005933656939305365
Iteration: 15, loss: 0.0005910006002523005
Iteration: 16, loss: 0.0005886448197998106
Iteration: 17, loss: 0.0005862982361577451
Iteration: 18, loss: 0.0005839611985720694
Iteration: 19, loss: 0.0005816334742121398
Iteration: 20, loss: 0.0005793151794932783
Iteration: 21, loss: 0.0005770060233771801
Iteration: 22, loss: 0.0005747061804868281
Iteration: 23, loss: 0.0005724154179915786
Iteration: 24, loss: 0.0005701337940990925
Iteration: 25, loss: 0.0005678612506017089
Iteration: 26, loss: 0.0005655977874994278
Iteration: 27, loss: 0.0005633432883769274
Iteration: 28, loss: 0.0005610977532342076
Iteration: 29, loss: 0.0005588610656559467
Iteration: 30, loss: 0.0005566334584727883
Iteration: 31, loss: 0.0005544146406464279
Iteration: 32, loss: 0.0005522047867998481
Iteration: 33, loss: 0.0005500036641024053
Iteration: 34, loss: 0.0005478113889694214
Iteration: 35, loss: 0.0005456277285702527
Iteration: 36, loss: 0.0005434528575278819
Iteration: 37, loss: 0.0005412864848040044
Iteration: 38, loss: 0.0005391290178522468
Iteration: 39, loss: 0.0005369800492189825
Iteration: 40, loss: 0.0005348395206965506
Iteration: 41, loss: 0.0005327076651155949
Iteration: 42, loss: 0.0005305844824761152
Iteration: 43, loss: 0.0005284695071168244
Iteration: 44, loss: 0.0005263631464913487
Iteration: 45, loss: 0.0005242649931460619
Iteration: 46, loss: 0.0005221752799116075
Iteration: 47, loss: 0.0005200940067879856
Iteration: 48, loss: 0.0005180207663215697
Iteration: 49, loss: 0.0005159559659659863
Iteration: 50, loss: 0.0005138994310982525
Iteration: 51, loss: 0.0005118509288877249
Iteration: 52, loss: 0.000509810633957386
Iteration: 53, loss: 0.0005077786045148969
Iteration: 54, loss: 0.0005057546077296138
Iteration: 55, loss: 0.0005037385853938758
Iteration: 56, loss: 0.0005017306539230049
Iteration: 57, loss: 0.00049973075510934
Iteration: 58, loss: 0.0004977386561222374
Iteration: 59, loss: 0.0004957547062076628
Iteration: 60, loss: 0.0004937786725349724

```
Iteration: 61, loss: 0.0004918103804811835
Iteration: 62, loss: 0.0004898500046692789
Iteration: 63, loss: 0.0004878975742030889
Iteration: 64, loss: 0.00048595271073281765
Iteration: 65, loss: 0.0004840158799197525
Iteration: 66, loss: 0.00048208649968728423
Iteration: 67, loss: 0.0004801649774890393
Iteration: 68, loss: 0.00047825108049437404
Iteration: 69, loss: 0.0004763447504956275
Iteration: 70, loss: 0.0004744461621157825
Iteration: 71, loss: 0.0004725549661088735
Iteration: 72, loss: 0.0004706713370978832
Iteration: 73, loss: 0.00046879539149813354
Iteration: 74, loss: 0.00046692678006365895
Iteration: 75, loss: 0.0004650656774174422
Iteration: 76, loss: 0.00046321196714416146
Iteration: 77, loss: 0.0004613656783476472
Iteration: 78, loss: 0.000459526723716408
Iteration: 79, loss: 0.00045769510325044394
Iteration: 80, loss: 0.00045587081694975495
Iteration: 81, loss: 0.00045405369019135833
Iteration: 82, loss: 0.00045224386849440634
Iteration: 83, loss: 0.00045044123544357717
Iteration: 84, loss: 0.00044864576193504035
Iteration: 85, loss: 0.000446857389761135
Iteration: 86, loss: 0.00044507632264867425
Iteration: 87, loss: 0.0004433022113516927
Iteration: 88, loss: 0.0004415352887008339
Iteration: 89, loss: 0.00043977529276162386
Iteration: 90, loss: 0.0004380222235340625
Iteration: 91, loss: 0.0004362761974334717
Iteration: 92, loss: 0.00043453715625219047
Iteration: 93, loss: 0.00043280518730171025
Iteration: 94, loss: 0.00043107994133606553
Iteration: 95, loss: 0.00042936153477057815
Iteration: 96, loss: 0.0004276501713320613
Iteration: 97, loss: 0.00042594553087837994
Iteration: 98, loss: 0.0004242476134095341
Iteration: 99, loss: 0.00042255656444467604
Iteration: 100, loss: 0.0004208721511531621
```

```
In [21]: plt.plot(losses)
plt.xlabel('Epoca')
plt.ylabel('Cost')
plt.show()
```

In []: